

Accelerator-Based Programming - GPU programming with Kokkos

Jörn Zimmerling

September 26, 2022



UPPSALA
UNIVERSITET

Programming of accelerators

- Many big research codes target a variety of platforms
 - Classical CPU-based systems (Intel, AMD, arm)
 - NVIDIA GPUs
 - AMD GPUs
 - Intel GPUs
 - ...
- How to address difference in architectures?
 - Separately implementing everything on multiple architectures not realistic
 - Write and maintain code with multiple paths - `#ifdef`
 - Needs highly skilled people on several architectures
 - Separately implementing important algorithms on multiple architectures works sometimes
 - Needs big enough team to drive it
 - Separation of concerns important: Multiple “Back-ends” (CPU,GPU) for key algorithms, but abstract it away from application code
- Create abstraction layer that targets generic loops



Kokkos

- Targets a single-source implementation with C++ programming language
- Create execution policy to important computational kernels (for loop, reduction, . . .) and express work as computational body (code to perform unit of work)
- Descriptive programming model
- Compile the code for different targets: CPU/GPU
- Select different data layouts adapted to hardware architecture
 - Map algorithm parallelism to hardware parallelism
 - Granularity of data access from different threads
- Goals of Kokkos:
 - Ease of use
 - Flexibility also to large codes
 - Performance

→ Performance portability



Kokkos in a nutshell

- Model for data parallelism
 - Use **parallel patterns** and **execution policies** to execute **compute bodies**
 - Example 1: parallel loops with the `parallel_for` pattern (e.g. vector add)

```
parallel_for ( "stream_triad" , N , [=] (int i) {  
    z( i) = a * x( i) + y( i) ;  
}) ;
```

- [=] - lambda capture list ([=] value, [&] reference) [c++ 11]
- Example 2: Reductions combine results from loop iterations

```
float result;  
parallel_reduce ( " summation " , N , [=] ( int i, float &lres ) {  
    lres += x( i) ;  
}, result) ;
```

- Flexibility choose data-layout problem with multi-dimensional array abstraction
- Execution and memory spaces to control
 - Where data lives
 - Where code executes



Learning Kokkos

- We will make use of extensive tutorial material provided by
- Kokkos team
 - General Kokkos GitHub project organization
 - <https://github.com/kokkos>
 - Lecture material: <https://github.com/kokkos/kokkos-tutorials/tree/main/LectureSeries>
- We will discuss
 - Lecture 1: Introduction
 - Lecture 2: Memory views and spaces
 - Lecture 3: Multi-dimensional loops and data structures
- Get Kokkos from <https://github.com/kokkos/kokkos>
- Lab consist out of tutorial's from Kokkos

