



# UPPSALA UNIVERSITET

Accelerator-Based Programming - 1TD055

ASSIGNMETN 4: KOKKOS

Jyong-Jhih Lin

October 10, 2022

## 0 Hardware information

snowy CPU:

[illegible]

snowy memory:

```

1 Handle 0x1100, DMI type 17, 40 bytes
2 Memory Device
3 Array Handle: 0x1000
4 Error Information Handle: Not Provided
5 Total Width: 72 bits
6 Data Width: 64 bits
7 Size: 32 GB
8 Form Factor: DIMM
9 Set: None
10 Locator: PROC 1 DIMM 1
11 Bank Locator: Not Specified
12 Type: DDR3
13 Type Detail: Synchronous LRDIMM
14 Speed: 1333 MT/s
15 Manufacturer: HP
16 Serial Number: Not Specified
17 Asset Tag: Not Specified
18 Part Number: 647654-081
19 Rank: 4
20 Configured Memory Speed: 1333 MT/s
21 Minimum Voltage: 1.35 V
22 Maximum Voltage: 1.5 V
23 Configured Voltage: 1.35 V
24
25 #####
26 DDR3-1333 4-channel memory total bandwidth = 1333e6(T/s) * 64(bits) * 4(channels) / 8e9(GBytes/s)
27 = 42.656(GBytes/s)

```

nvidia T4, nvidia-smi:

```

1 | +-----+-----+-----+-----+-----+-----+-----+-----+
2 | | NVIDIA-SMI 515.65.01      Driver Version: 515.65.01    CUDA Version: 11.7   |
3 | +-----+-----+-----+-----+-----+-----+-----+-----+
4 | | GPU Name Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
5 | | Fan  Temp  Perf    Pwr:Usage/Cap|     Memory-Usage | GPU-Util  Compute M. |
6 | |                               |                  MIG M. |                       |
7 | +-----+-----+-----+-----+-----+-----+-----+-----+
8 | |    0   Tesla T4               On          00000000:08:00:0 Off |                    0  |
9 | | N/A   30C   P8              14W /  70W           2MiB / 15360MiB |          0%      Default |
10 | |                               |                  N/A   |                       |
11 | +-----+-----+-----+-----+-----+-----+-----+-----+
12 |
13 | +-----+-----+-----+-----+-----+-----+-----+-----+
14 | | Processes:                                |
15 | |   GPU   GI    CI        PID   Type   Process name                      GPU Memory |
16 | |       ID   ID                                 name                        Usage     |
17 | +-----+-----+-----+-----+-----+-----+-----+-----+

```

```

18 | No running processes found |
19 +-----+

```

#### nvidia T4, *deviceQuery*:

```

1 /sw/EasyBuild/snowy/software/CUDA/10.1.243-iccifort-2019.5.281/extras/demo_suite/deviceQuery Starting...
2
3 CUDA Device Query (Runtime API) version (CUDA static linking)
4
5 Detected 1 CUDA Capable device(s)
6
7 Device 0: "Tesla T4"
8   CUDA Driver Version / Runtime Version      11.7 / 10.1
9   CUDA Capability Major/Minor version number: 7.5
10  Total amount of global memory:              14972 MBytes (15699148800 bytes)
11  (40) Multiprocessors, ( 64) CUDA Cores/MP:  2560 CUDA Cores
12  GPU Max Clock rate:                        1590 MHz (1.59 GHz)
13  Memory Clock rate:                         5001 Mhz
14  Memory Bus Width:                          256-bit
15  L2 Cache Size:                             4194304 bytes
16  Maximum Texture Dimension Size (x,y,z)      1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
17  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
18  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
19  Total amount of constant memory:             65536 bytes
20  Total amount of shared memory per block:     49152 bytes
21  Total number of registers available per block: 65536
22  Warp size:                                   32
23  Maximum number of threads per multiprocessor: 1024
24  Maximum number of threads per block:         1024
25  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
26  Max dimension size of a grid size (x,y,z):   (2147483647, 65535, 65535)
27  Maximum memory pitch:                       2147483647 bytes
28  Texture alignment:                          512 bytes
29  Concurrent copy and kernel execution:        Yes with 3 copy engine(s)
30  Run time limit on kernels:                   No
31  Integrated GPU sharing Host Memory:          No
32  Support host page-locked memory mapping:     Yes
33  Alignment requirement for Surfaces:         Yes
34  Device has ECC support:                      Enabled
35  Device supports Unified Addressing (UVA):    Yes
36  Device supports Compute Preemption:         Yes
37  Supports Cooperative Kernel Launch:         Yes
38  Supports MultiDevice Co-op Kernel Launch:   Yes
39  Device PCI Domain ID / Bus ID / location ID: 0 / 8 / 0
40  Compute Mode:
41     < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
42
43 deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 11.7, CUDA Runtime Version = 10.1, NumDevs = 1,
44   Device0 = Tesla T4
45 Result = PASS

```

#### DELL Precision 7760 CPU:

```

1 Architecture:                x86_64
2 CPU op-mode(s):              32-bit, 64-bit
3 Byte Order:                  Little Endian
4 Address sizes:               39 bits physical, 48 bits virtual
5 CPU(s):                      12
6 On-line CPU(s) list:        0-11
7 Thread(s) per core:         2
8 Core(s) per socket:         6
9 Socket(s):                   1
10 NUMA node(s):               1
11 Vendor ID:                   GenuineIntel
12 CPU family:                  6
13 Model:                      141
14 Model name:                  Intel(R) Xeon(R) W-11855M CPU @ 3.20GHz
15 Stepping:                    1
16 CPU MHz:                     3200.000
17 CPU max MHz:                 4900.0000
18 CPU min MHz:                 800.0000
19 Bogomips:                    6374.40
20 Virtualization:              VT-x
21 L1d cache:                   288 KiB
22 L1i cache:                   192 KiB
23 L2 cache:                    7.5 MiB
24 L3 cache:                    18 MiB
25 NUMA node0 CPU(s):          0-11
26 Vulnerability Itlb multihit: Not affected
27 Vulnerability L1tf:          Not affected
28 Vulnerability Mds:           Not affected
29 Vulnerability Meltdown:      Not affected
30 Vulnerability Mmio stale data: Not affected
31 Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl and seccomp
32 Vulnerability Spectre v1:     Mitigation; usercopy/swapgs barriers and __user pointer sanitization
33 Vulnerability Spectre v2:     Mitigation; Enhanced IBRS, IBPB conditional, RSB filling
34 Vulnerability Srbds:          Not affected

```

```

35 Vulnerability Tsx async abort:    Not affected
36 Flags:                            fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36
    clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art
    arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc cpuid aperfperf tsc_known_freq pni pclmulqdq
    dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe
    popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb cat_l2
    invpcid_single cdp_l2 ssbd ibrs ibpb stibp ibrs_enhanced tpr_shadow vnmi flexpriority ept vpid ept_ad
    fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid rdt_a avx512f avx512dq rdseed adx smap avx512ifma
    clflushopt clwb intel_pt avx512cd sha_ni avx512bw avx512vl xsaveopt xsavec xgetbv1 xsaves
    split_lock_detect dtherm ida arat pln pts hwp hwp_notify hwp_act_window hwp_epp hwp_pkg_req avx512vbmi
    umip pku ospke avx512_vbmi2 gfni vaes vpclmulqdq avx512_vnni avx512_bitalg tme avx512_vpopcntdq rdpid
    movdiri movdir64b fsrm avx512_vp2intersect md_clear flush_l1d arch_capabilities

```

#### DELL Precision 7760 memory:

```

1 Handle 0x1100, DMI type 17, 92 bytes
2 Memory Device
3   Array Handle: 0x1000
4   Error Information Handle: Not Provided
5   Total Width: 72 bits
6   Data Width: 64 bits
7   Size: 32 GB
8   Form Factor: SODIMM
9   Set: None
10  Locator: DIMM C
11  Bank Locator: BANK 0
12  Type: DDR4
13  Type Detail: Synchronous
14  Speed: 2933 MT/s
15  Manufacturer: 01980000802C
16  Serial Number: 97B0B609
17  Asset Tag: 04212100
18  Part Number: 9965657-029.A00G
19  Rank: 2
20  Configured Memory Speed: 2933 MT/s
21  Minimum Voltage: Unknown
22  Maximum Voltage: Unknown
23  Configured Voltage: 1.2 V
24  Memory Technology: DRAM
25  Memory Operating Mode Capability: Volatile memory
26  Firmware Version: Not Specified
27  Module Manufacturer ID: Bank 2, Hex 0x98
28  Module Product ID: Unknown
29  Memory Subsystem Controller Manufacturer ID: Unknown
30  Memory Subsystem Controller Product ID: Unknown
31  Non-Volatile Size: None
32  Volatile Size: 32 GB
33  Cache Size: None
34  Logical Size: None
35
36 #####
37 DDR4-2933 2-channel memory total bandwidth = 2933e6(T/s) * 64(bits) * 2(channels) / 8e9(GBytes/s)
38 = 46.928(GBytes/s)

```

#### nvidia A3000, nvidia-smi:

```

1 +-----+
2 | NVIDIA-SMI 510.47.03      Driver Version: 510.47.03      CUDA Version: 11.6      |
3 |-----+-----+-----+-----+-----+-----+
4 | GPU   Name               Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
5 | Fan   Temp   Perf         Pwr:Usage/Cap|               Memory-Usage | GPU-Util  Compute M. |
6 |-----+-----+-----+-----+-----+-----+-----+
7 | 0     0    NVIDIA RTX A300... Off      | 00000000:01:00:0  On    |                     N/A |
8 | N/A    58C     P0      36W /  N/A     | 1606MiB /  6144MiB |    100%      Default  |
9 |-----+-----+-----+-----+-----+-----+
10 |
11 +-----+
12
13 +-----+
14 | Processes:
15 | GPU   GI    CI          PID    Type   Process name                      GPU Memory
16 |    ID   ID                                 Usage
17 |-----+-----+-----+-----+-----+-----+
18 | 0     N/A  N/A       3964      G   /usr/lib/xorg/Xorg                  109MiB
19 | 0     N/A  N/A      12824      G   /usr/lib/xorg/Xorg                  603MiB
20 | 0     N/A  N/A      12940      G   /usr/bin/gnome-shell                271MiB
21 | 0     N/A  N/A      30513      G   ...308337019390783085,131072        481MiB
22 | 0     N/A  N/A      215156     G   ...R2021a/bin/glnxa64/MATLAB         3MiB
23 | 0     N/A  N/A      216972     G   ...GL_KHR_blend_equation_adv         5MiB
24 | 0     N/A  N/A      284196     C   ./stream_triad_cuda                  113MiB
25 +-----+

```

#### nvidia A3000, deviceQuery:

```

1 ./deviceQuery Starting...

```

```

2
3  CUDA Device Query (Runtime API) version (CUDA static linking)
4
5  Detected 1 CUDA Capable device(s)
6
7  Device 0: "NVIDIA RTX A3000 Laptop GPU"
8  CUDA Driver Version / Runtime Version      11.6 / 11.6
9  CUDA Capability Major/Minor version number: 8.6
10 Total amount of global memory:              5913 MBytes (6200098816 bytes)
11 (32) Multiprocessors, (128) CUDA Cores/MP: 4096 CUDA Cores
12 GPU Max Clock rate:                        1560 MHz (1.56 GHz)
13 Memory Clock rate:                         5501 Mhz
14 Memory Bus Width:                          192-bit
15 L2 Cache Size:                             3145728 bytes
16 Maximum Texture Dimension Size (x,y,z)      1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
17 Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
18 Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
19 Total amount of constant memory:             65536 bytes
20 Total amount of shared memory per block:     49152 bytes
21 Total number of registers available per block: 65536
22 Warp size:                                  32
23 Maximum number of threads per multiprocessor: 1536
24 Maximum number of threads per block:         1024
25 Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
26 Max dimension size of a grid size (x,y,z):   (2147483647, 65535, 65535)
27 Maximum memory pitch:                       2147483647 bytes
28 Texture alignment:                          512 bytes
29 Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
30 Run time limit on kernels:                   Yes
31 Integrated GPU sharing Host Memory:          No
32 Support host page-locked memory mapping:     Yes
33 Alignment requirement for Surfaces:          Yes
34 Device has ECC support:                      Disabled
35 Device supports Unified Addressing (UVA):    Yes
36 Device supports Compute Preemption:         Yes
37 Supports Cooperative Kernel Launch:         Yes
38 Supports MultiDevice Co-op Kernel Launch:   Yes
39 Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
40 Compute Mode:
41     < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
42
43 deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 11.6, CUDA Runtime Version = 11.6, NumDevs = 1,
44 Device0 = NVIDIA RTX A3000 Laptop GPU
45 Result = PASS

```

## 1 Task 1

### 1.1 a

Listing 1: explicitly data transfer and initiation

```

34 Kokkos::View<fptype*[3][3], Kokkos::memlayout, MemSpace> J( "J", n_element);
35 Kokkos::View<fptype*[4][4], Kokkos::memlayout, MemSpace> A( "A", n_element);
36
37 Kokkos::View<fptype*[3][3], Kokkos::memlayout, MemSpace>::HostMirror h_J = Kokkos::create_mirror_view( J
38 );
39 Kokkos::View<fptype*[4][4], Kokkos::memlayout, MemSpace>::HostMirror h_A = Kokkos::create_mirror_view( A
40 );
41
42 Kokkos::Timer timer;
43 for ( int i =0; i<n_element; i++){
44     h_J(i,0,0) = 3;
45     h_J(i,0,1) = 1;
46     h_J(i,0,2) = 1;
47     h_J(i,1,0) = 1;
48     h_J(i,1,1) = 3;
49     h_J(i,1,2) = 1;
50     h_J(i,2,0) = 1;
51     h_J(i,2,1) = 1;
52     h_J(i,2,2) = 3;
53 }
54 double t_j_init = timer.seconds();
55
56 timer.reset();
57 Kokkos::deep_copy( J, h_J );
58 Kokkos::fence();
59 double t_j_init_copy = timer.seconds();
60 timer.reset();
61 Kokkos::deep_copy( A, h_A );
62 Kokkos::fence();
63 double t_a_init_copy = timer.seconds();

```

Listing 2: explicitly data transfer and initiation

```

103 timer.reset();
104 Kokkos::deep_copy( h_J, J );
105 Kokkos::fence();
106 double t_j_end_copy = timer.seconds();
107 timer.reset();
108 Kokkos::deep_copy( h_A, A );
109 Kokkos::fence();
110 double t_a_end_copy = timer.seconds();

```

## 1.2 b

Listing 3: implementations both with float and double numbers

```

14 #define fptype float

```

By changing the “fptype” definition, you can change between “float” and “double” type.

## 1.3 c

Too many places to be listed out. Please check the source code directly.

## 1.4 d

Listing 4: implementations both with float and double numbers

```

15 #define memlayout LayoutLeft

```

By changing the “memlayout” definition, you can change between “LayoutLeft” and “LayoutRight”.

## 2 Task 2




Figure 1: GPU=A3000, CPU=W-11855M

execution environment configuration:

```
1 OMP_NUM_THREADS=6 OMP_PROC_BIND=spread OMP_PLACES=cores ./ABP_lab4_t1.host
```

In the DRAM region, the CPU performance achieves the DDR4-2933 1-channel bandwidth as expected and is consistent with the previous lab experiments. The GPU performance is also close to the GDDR6 peak bandwidth. The CPU and GPU are both in memory bandwidth bound when the matrix size is in the DRAM region. Because they are in the memory bandwidth bound, the double and float type version have the same performance bandwidth.

The CPU version achieves its maximum performance at  $N$  around  $10^5$  which is in the cache region. When the matrix size is in the cache region, the cache memory bandwidth is no longer an issue and it becomes a compute bound. Therefore, the double type version has twice bandwidth performance as the float type one.

For the GPU version, because the cache size per thread is relatively much smaller than the CPU, the cache bandwidth only matters when the  $N$  is very small. Therefore, the GPU-double and GPU-float version basically have the same performance after  $N=10^4$ .

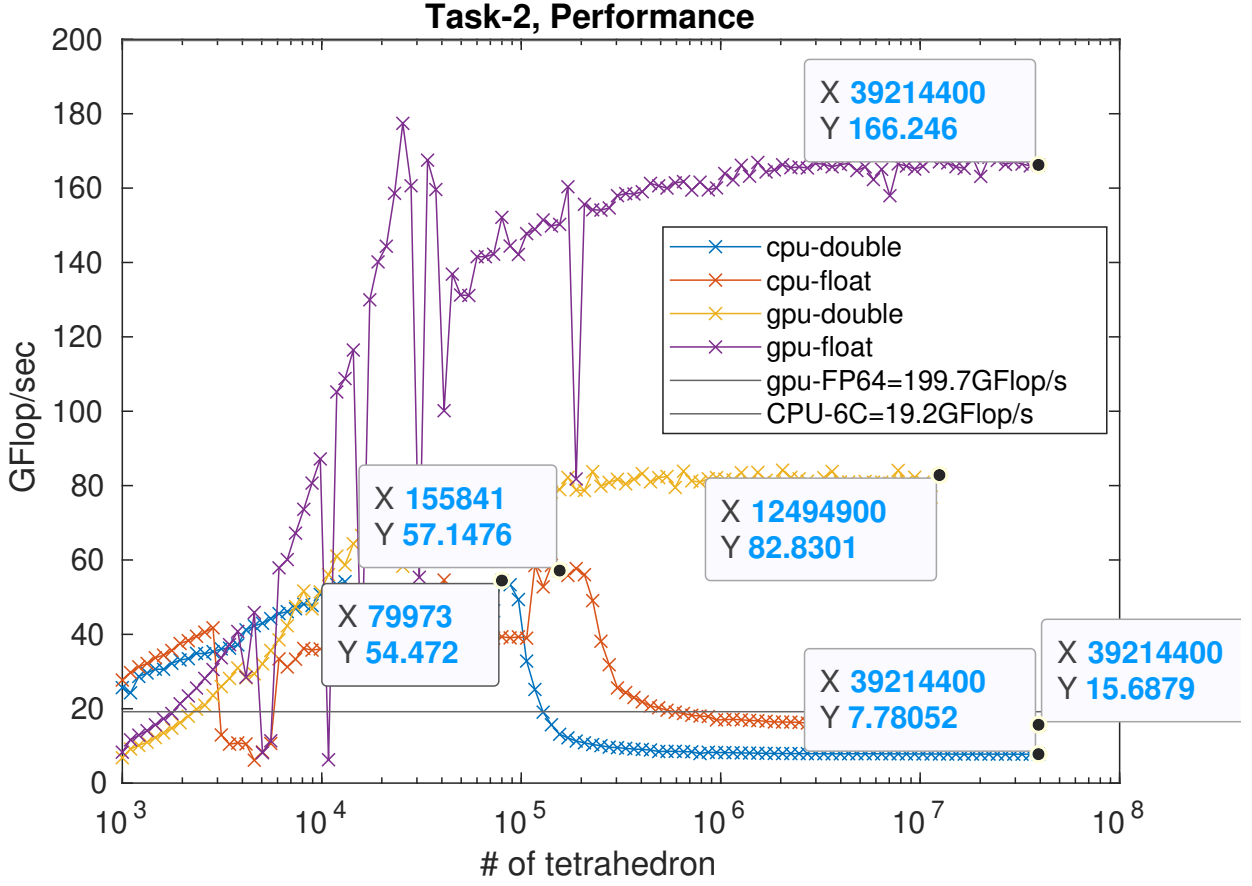


Figure 2: GPU=A3000, CPU=W-11855M

In the DRAM region, the CPU-double version has half of the compute performance of the CPU-float version because of the memory bandwidth bound. In the cache region, the CPU-double and CPU-float version have the same compute performance because the cache memory bandwidth is no longer an issue.

For the GPU version, the GPU-float version has twice compute performance as the GPU-double version because of the memory bandwidth bound, except when the  $N$  is very small and in the cache region.

### 3 Task 3

According to the hardware specification, <https://www.techpowerup.com/gpu-specs/rtx-a3000-mobile.c3806>, the peak compute performance and the peak memory bandwidth for NVIDIA RTX A3000 GPU are 12.78 TFLOPS for FP32, 199.7 GFLOPS for FP64, and 264.0 GB/s.

The CPU performance could be estimated by the following formula:

$$\frac{P}{GFlop/s} = \frac{clock\ freq.}{GHz} * (\# of\ cores) * (SIMD\ width) * \frac{instructions}{cycle} * \frac{FLOPs}{instruction}$$

For the Intel Xeon W-11855M Processor, the clock frequency is dynamic and I will use its base frequency 3.2GHz as an estimation. It has 6 physical cores. In this program, I assume that it does not use any SIMD instructions, even with the optimization O3. Usually, the instruction per cycle is 2. I assume that the compiler does not use any FMA instructions, even with the optimization O3. Therefore, the FLOPs per instruction is 1. The final estimation of the CPU performance for the Kokkos program is 38.4 GFlop/s

For the maximum peak compute performance, the SIMD width should be 512/32 and the FLOPs per instruction should be 2 for FMA. The estimation would be 1228.8 GFlop/s.



If the compute algorithm does not change, the limiting resource for the CPU and GPU version are the memory bandwidth.

A practical approach to improve the performance is to have a better memory layout and access pattern. It would help utilize the memory prefetch, reduce the memory bank conflict, and the cache line false sharing. These techniques work for both CPU and GPU.

## 4 Task 4

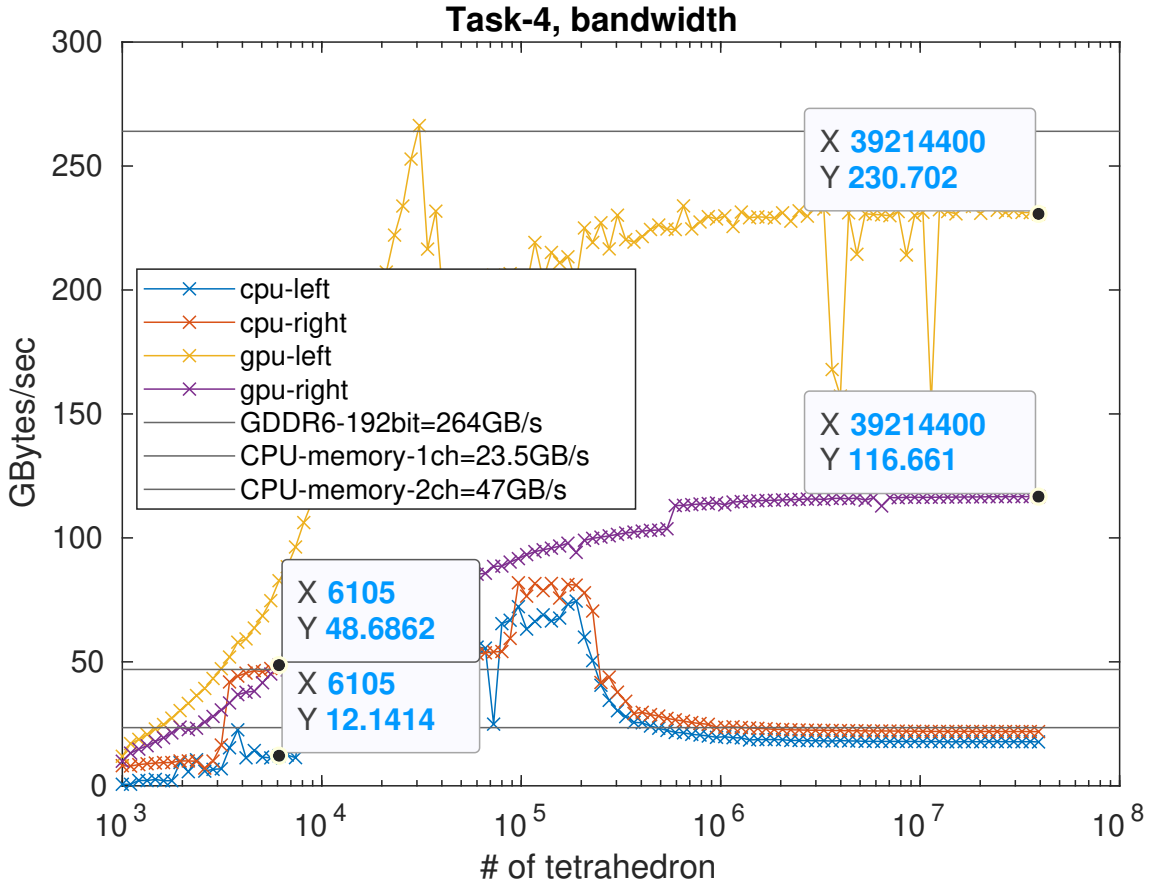


Figure 3: GPU=A3000, CPU=W-11855M, float




Figure 4: GPU=A3000, CPU=W-11855M, float

For the GPU, the optimal memory layout should be LayoutLeft in all memory region. The wrong memory layout could reduce the performance as much as half.

For the CPU, the optimal memory layout should be LayoutRight, especially in the cache memory region. When the matrix size is large and in the DRAM region, the performance difference between LayoutLeft and LayoutRight is relatively small. I guess the performance difference is reduced by the “random-access” feature of the memory and the CPU cache as a buffer. The performance difference in the cache memory region is large because of the “cache line false sharing” effect. It is an important issue for the OpenMP parallelization.

## 5 Task 5

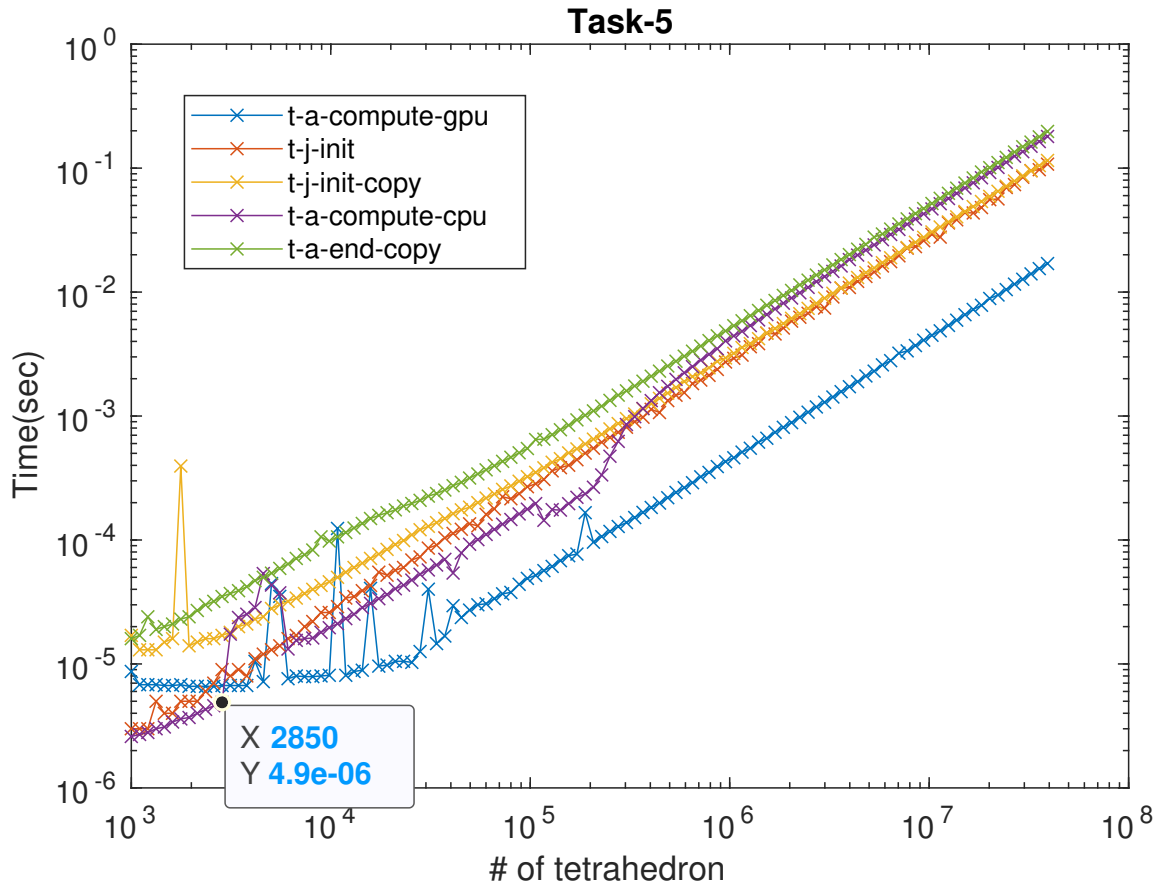


Figure 5: GPU=A3000, CPU=W-11855M, float

## 6 Task 6

First of all, for a specific problem size  $N$ , the program should be executed on CPU or GPU should be decided by the over all execution time, which includes the initialization, data transfer, and execution on all devices. However, the task 6 only asks to compare a specific section of the program, which does not always reflect the over all execution time.

- the Jacobians are computed from the vertices and connectivity information  
For this section, because the “t-j-init” is always smaller than the “t-j-init-copy”, it would be more efficient to calculate the J matrix on CPU than GPU.
- the element matrices are computed as above  
For this section, because the “t-a-compute-gpu” is smaller than the “t-j-compute-cpu”, it would be more efficient to calculate the A matrix on GPU than CPU. However, when the  $N < 2850$ , the GPU computation is less efficient than the CPU. It means that for a small size calculation, it would be better to just use the CPU.
- the element matrices are assembled into a sparse matrix  
I do not understand this question. I cannot answer it.
- the linear system is solved with an iterative solver  
For the computation heavy section, this lab already shows that in general the GPU could have a better compute performance. Therefore, it should be executed on GPU.